

TUTORIAL: INTRODUCTION TO VALGRIND

RODRIGO RANDEL
RODRIGO.RANDEL@GERAD.CA



**POLYTECHNIQUE
MONTRÉAL**

LE GÉNIE
EN PREMIÈRE CLASSE

REFERENCE

This presentations was highly inspired by [Valgrind Documentation](#)

REPOSITORY WITH EXAMPLES

https://github.com/rodrigorandel/introduction_to_valgrind

WHAT IS CODE DEBUGGING AND PROFILING?

MOTIVATION

- ▶ After **developing your code**, you want make sure that:
 - ▶ It is free of bugs
 - ▶ Faster as possible

VALGRIND TOOLS THAT CAN AUTOMATICALLY DETECT MANY MEMORY MANAGEMENT AND THREADING BUGS, AND PROFILE YOUR PROGRAMS IN DETAIL

AN OVERVIEW OF VALGRIND

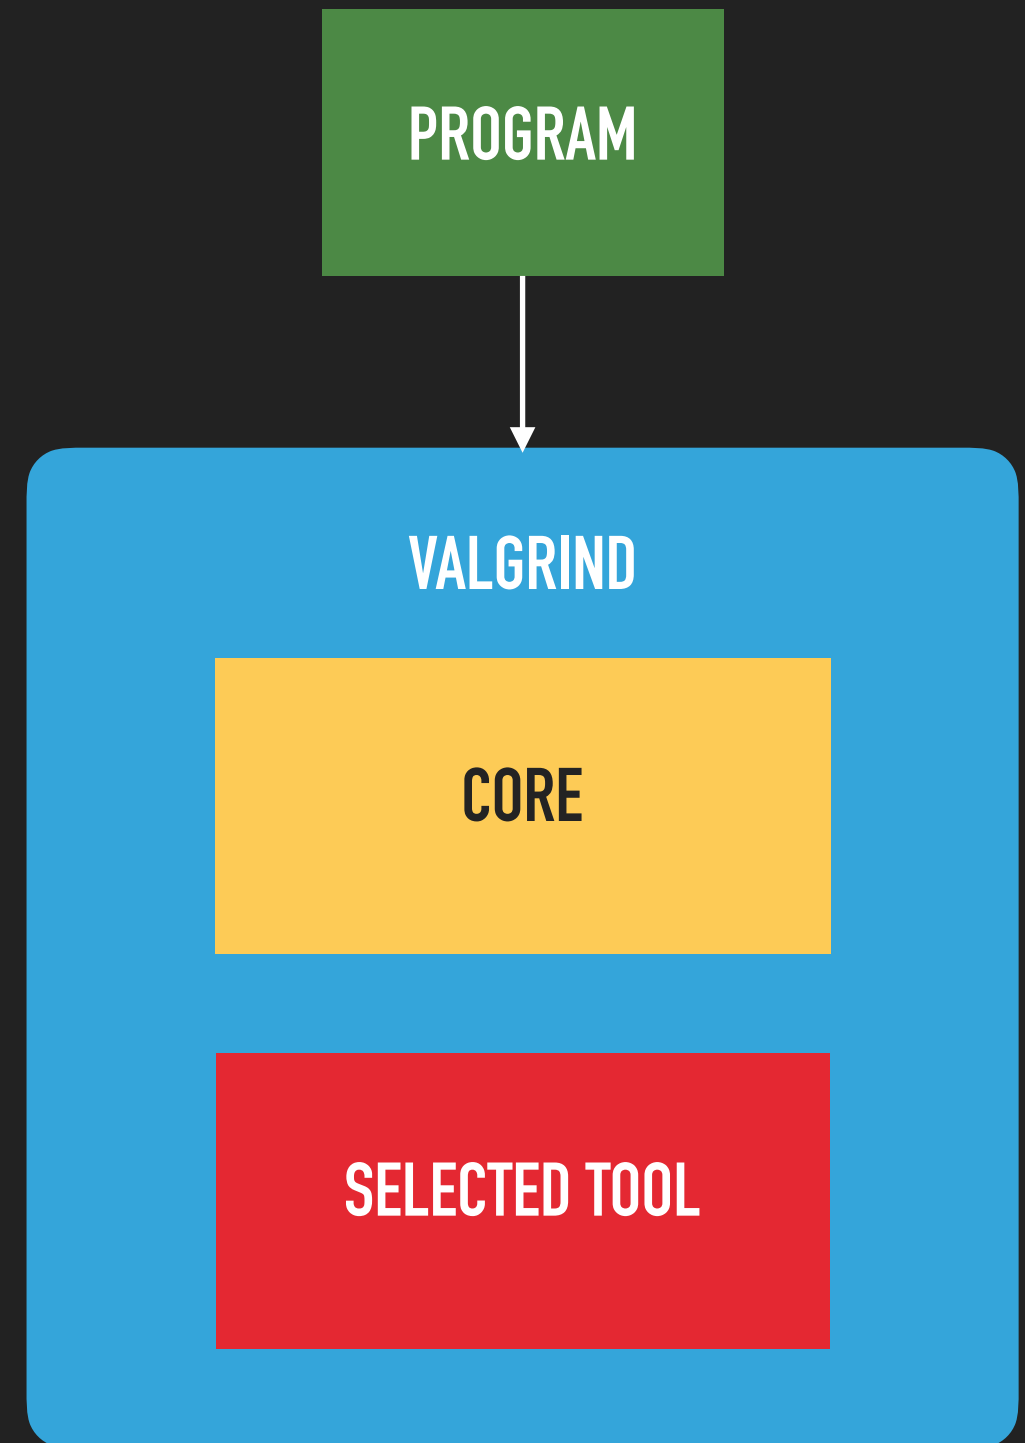
- ▶ Valgrind is an **instrumentation framework** for building dynamic **analysis tools**.
- ▶ Valgrind's architecture is **modular**, so new tools can be created easily
- ▶ Valgrind is designed to be as non-intrusive as possible. It works directly with **existing executables**

AN OVERVIEW OF VALGRIND

- ▶ The Valgrind distribution includes the following debugging and profiling tools:
 - ▶ **Memcheck** is a memory error detector. It helps you make your **programs more correct** (particularly those written in C and C++).
 - ▶ **Callgrind** is a call-graph generating cache profiler. It helps **identifying the bottleneck** of your program.
 - ▶ **Massif** is a heap profiler. It helps you make your programs **use less memory**.
 - ▶ **Cachegrind** is a cache and branch-prediction profiler. It helps you make your programs **run faster**.
 - ▶ **Helgrind** is a thread error detector. It helps you make your **multi-threaded programs more correct**.
 - ▶ **DRD** is also a thread error detector. It is similar to Helgrind but uses different analysis techniques and so may find different problems.

HOW VALGRIND WORKS

- ▶ Your **program** runs on a synthetic CPU provided by the **Valgrind core**:
 - ▶ When the code is executed for the first time, the core hands the code to the **selected tool**.
 - ▶ The tool adds its own instrumentation code
 - ▶ Hands the result back to the core
 - ▶ The core coordinates the continued execution of this instrumented code.



HOW VALGRIND WORKS

- ▶ No need to recompile, relink, or otherwise modify the program to be checked

```
▶ valgrind [valgrind-options] your-prog [your-prog-options]
```

- ▶ The most important option is **--tool** which dictates which Valgrind tool to run.

```
▶ valgrind --tool=memcheck ls -l
```


UNDERSTANDING VALGRIND'S OUTPUT (MEMCHECK)

```
rodrigorandel@valgrindtutorial:~/Examples$ valgrind --tool=memcheck ./ex1
==2515== Memcheck, a memory error detector
==2515== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==2515== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright
info
==2515== Command: ./ex1
==2515==
==2515== Invalid write of size 4
==2515==    at 0x1086B4: main (ex1.cpp:5)
==2515== Address 0x5b20ca8 is 0 bytes after a block of size 40 alloc'd
==2515==    at 0x4C3089F: operator new[](unsigned long) (in /usr/lib/
valgrind/vgpreload_memcheck-amd64-linux.so)
==2515==    by 0x10868B: main (ex1.cpp:3)
==2515==
==2515==
==2515== HEAP SUMMARY:
==2515==    in use at exit: 0 bytes in 0 blocks
==2515== total heap usage: 2 allocs, 2 frees, 72,744 bytes allocated
==2515==
==2515== All heap blocks were freed -- no leaks are possible
==2515==
==2515== For counts of detected and suppressed errors, rerun with: -v
==2515== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from 0)
```

PID

COMMENTARY

```
2 int main(){
3     int* arr = new int[10];
4     for( int i=0; i <= 10; i++){
5         arr[i] = i;
6     }
7
8     delete [] arr;
9     return 0;
10 }
```

COMMENTARY

VALGRIND CORE (BASIC) OPTIONS

- ▶ `--tool=<toolname>` [default: memcheck]
Run the Valgrind tool called toolname
- ▶ `-h --help`
Show help for all options, both for the core and for the selected tool.
- ▶ `-q --quiet`
Run silently, and only print error messages.
- ▶ `-v --verbose`
Gives extra information on various aspects of your program
- ▶ `--time-stamp=<yes|no>` [default: no]
Each message is preceded with an indication of the elapsed wallclock time since startup
- ▶ `--log-file=<filename>`
Specifies that Valgrind should send all of its messages to the specified file

DEMO 1: EX1.CPP

REVIEW FROM DEMO 1

▶ **When using Memcheck with C/C++**

▶ **Compile with debugging info (-g for C/C++).**

- ▶ Without debugging info, the best Valgrind tools will be able to do is guess which function a particular piece of code belongs to, which makes both error messages and profiling output nearly useless.

▶ **Be careful when compiling with -O1 (and above)**

- ▶ The best solution would be is to turn off optimisation altogether. Since this often makes things unmanageably slow, a reasonable compromise is to use -O1. On rare occasions, compiler optimisations (at -O2 and above, and sometimes -O1) have been observed to generate code which fools Memcheck.

SUPPRESSING ERRORS

SUPPRESSING ERRORS

- ▶ The error-checking tools detect numerous problems in the **system libraries**, such as:
 - ▶ C library, which come **preinstalled with your OS**
 - ▶ Third-party software
- ▶ You can't easily fix these, but you **don't want to see these errors** (and yes, there are many!)
- ▶ Valgrind reads a list of errors to suppress at startup - **default suppression file** is created during installation (**demo with -v**).

WRITING SUPPRESSING ERRORS (GUIDE)

- ▶ Asking valgrind to generate suppressions automatically.

--gen-suppressions=<yes|no|all> [default: no]

--suppressions=file

```
{  
    name_of_suppression  
    tool_name:supp_kind  
    (optional extra info for some suppression types)  
    caller0 name, or /name/of/so/file.so  
    caller1 name, or ditto  
    (optionally: caller2 name)  
    (optionally: caller3 name)  
}
```

TEMPLATE

WRITING SUPPRESSING ERRORS (GUIDE)

- ▶ Write your own suppression file:
 - ▶ Use the `--demangle=no` option to get the mangled names in your error messages

```
==6386== 40 bytes in 1 blocks are definitely lost in loss record 1 of 1
==6386==      at 0x4C3089F: operator new[](unsigned long) (in /usr/lib/valgrind/
vgpreload_memcheck-amd64-linux.so)
==6386==      by 0x10891B: test() (ex1.cpp:6)
==6386==      by 0x108957: main (ex1.cpp:13)
```



```
==6388== 40 bytes in 1 blocks are definitely lost in loss record 1 of 1
==6388==      at 0x4C3089F: Znam (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==6388==      by 0x10891B: Z4testv (ex1.cpp:6)
==6388==      by 0x108957: main (ex1.cpp:13)
```

CALLERS

DEMO 2: EX2.CPP

DEBUGGING YOUR PROGRAM USING VALGRIND

VALGRIND GDBSERVER AND GDB

- ▶ Recall: Your program runs in a **synthetic CPU** provided by Valgrind
 - ▶ A debugger **cannot debug** your program when it runs on Valgrind.
- ▶ **Valgrind gdbserver** provides a fully debuggable program under Valgrind
 - ▶ GDB also provides an interactive usage of **Valgrind core** or **tool functionalities**,

VALGRIND GDBSERVER (BASICS) OPTIONS

- ▶ `--vgdb=<no|yes|full>` [default: yes]

Provide "gdbserver" functionality when `--vgdb=yes` or `--vgdb=full` is specified

`--vgdb=full` provides more precise breakpoints (incurs performance overheads)

- ▶ `--vgdb-error=<number>` [default: 999999999]

Tools that report errors will wait for "number" errors to be reported before freezing the program and waiting for you to connect with GDB

```
▶ valgrind --vgdb=yes --vgdb-error=0 ./ex1
```

```
==1584== TO DEBUG THIS PROCESS USING GDB: start GDB like this
```

```
==1584==    /path/to/gdb ./ex1
```

```
==1584== and then give GDB the following command
```

```
==1584==    target remote | /usr/lib/valgrind/../../bin/vgdb --pid=1584
```

```
==1584== --pid is optional if only one valgrind process is running
```

DEBUGGING USING VALGRIND

```
▶ gdb ./ex1
Reading symbols from ./ex1...done.
(gdb) target remote | vgdb --pid=1584
Remote debugging using | vgdb --pid=1584
relaying data between gdb and process 1584
0x000000004000ea0 in _start () from /lib64/ld-linux-x86-64.so
(gdb)
```

- ▶ **Breakpoints** can be inserted or deleted.
- ▶ Variables and register **values can be examined** or modified.
- ▶ **Execution can be controlled** (continue, step, next, stepi, etc).

VALGRIND MONITOR COMMANDS (COMPLETE LIST)

- ▶ (gdb) monitor v.info all_errors
shows all errors found so far.
- ▶ (gdb) monitor v.info last_error
shows the last error found.
- ▶ (gdb) monitor v.info n_errs_found
shows the number of errors found so far (current value of the --vgdb-error).
- ▶ (gdb) monitor v.set {gdb_output | log_output | mixed_output}
allows redirection of the Valgrind output (e.g. the errors detected by the tool).
- ▶ (gdb) monitor v.kill
requests the gdbserver to kill the process.

BASICS MEMCHECK MONITOR COMMANDS (COMPLETE LIST)

- ▶ (gdb) monitor leak_check [full*| summary| xtleak]
performs a leak check
- ▶ (gdb) monitor who_points_at <addr> [<len>]
shows all the locations where a pointer to addr is found

DEMO 3: EX3.CPP

MEMCHECK

MEMORY ERROR DETECTOR

- ▶ It can detect common problems in C and C++ programs.
 - ▶ Accessing memory that you **should not access**
 - ▶ Using **undefined values**, i.e. values that have not been initialized
 - ▶ **Incorrect freeing** of heap memory
 - ▶ **Memory leaks.**

PROBLEMS LIKE THESE CAN BE DIFFICULT TO FIND BY OTHER MEANS, OFTEN REMAINING UNDETECTED FOR LONG PERIODS, THEN CAUSING OCCASIONAL, DIFFICULT-TO-DIAGNOSE CRASHES.

COMMON ERROR MESSAGES

▶ Illegal read / Illegal write errors

- ▶ This happens when your program reads or writes memory at a place it shouldn't
- ▶ **Hint:** `--read-var-info` will run more slowly but may give a more detailed description of any illegal address

▶ Uninitialized values

- ▶ An uninitialized-value use error is reported when your program uses a value that is undefined
- ▶ **Hint:** `--track-origins=yes` easier to track down the root causes of uninitialized value errors

COMMON ERROR MESSAGES

► Memory leak detection

LEAK CASES

Pointer Chain	A Leak case	B Leak Case
R----->B		Directly Reachable
R---->A---->B	Directly Reachable	Indirectly Reachable
R B		Directly Lost
R A---->B	Directly Lost	Indirectly Lost
R-----?----->B		Possibly Lost
R---->A--?-->B	Directly Reachable	Possibly Lost

LEAK KINDS RETURNED BY MEMCHECK

Still Reachable	Definitely Lost	Indirectly Lost	Possibly Lost
-----------------	-----------------	-----------------	---------------

- **Hint:** `--leak-check=full` will give details for each definitely lost or possibly lost block, including where it was allocated

DEMO 4: EX4.CPP

MASSIF


OVERVIEW

- ▶ Measures **how much heap memory** your program uses
- ▶ Can help you **reduce the amount of memory** your program uses.
 - ▶ It can **speed up your program**
 - ▶ Reduces the chance that it **exhausts your machine's swap space.**
- ▶ Detailed information that indicates **which parts** of your program are responsible for **allocating the heap memory.**

BASIC USAGE

```
▶ valgrind --tool=massif your-prog [your-prog-options]
```

- ▶ After program termination, a profile data file named **massif.out.<pid>** is generated
- ▶ Run **ms_print** to see the information gathered;
- ▶ (or use a visualization plugin such as eclipse)



DEMO 5

(USING ECLIPSE)

CALLGRIND

OVERVIEW

- ▶ Profiling tool that records the **call history among functions** in a program's run as a call-graph.
- ▶ Collects flat profile data: **event counts** (data reads, cache misses, etc.) are attributed directly to the **function they occurred in**.
- ▶ Allows you to find the specific call chains starting from 'main' in which the **majority of the program's costs occur**.

BASIC USAGE

```
▶ valgrind --tool=callgrind your-prog [your-prog-options]
```

- ▶ After program termination, a profile data file named `callgrind.out.<pid>` is generated.
 - ▶ Use `callgrind_annotate callgrind.out.<pid>`
 - ▶ **Hint:** `--inclusive=yes`: Instead of using exclusive cost of functions as sorting order, use and show inclusive cost.
 - ▶ **Hint:** `--auto=yes` to get annotated source code for all relevant functions for which the source can be found
- ▶ **Graphical visualization** of the data with [KCachegrind](#)

DEMO 6

(USING KCACHEGRIND)

CACHEGRIND

OVERVIEW

- ▶ **Simulates** how your program interacts with a **machine's cache hierarchy**
 - ▶ I1, D1 and LL (last-level) caches.
- ▶ Gather statistics of **caches reads and misses**
- ▶ These **statistics** are presented for the **entire program** and for **each function** in the program.
- ▶ Can also annotate each line of source code in the program with the **counts that were caused directly by it**.

BASIC USAGE

```
▶ valgrind --tool=cachegrind your-prog [your-prog-options]
```

- ▶ After program termination, a profile data file named `cachegrind.out.<pid>` is generated.
 - ▶ Use `cg_annotate cachegrind.out.<pid>`
 - ▶ Use `cg_diff first_file second_file` to compare two profiling
- ▶ **Ideally used with optimizations flags**

DEMO 7

(USING ECLIPSE)

HELGRIND & DRD

THREAD ERROR DETECTORS:

- ▶ Helgrind (Eclipse support)

- ▶ Detect synchronization errors in C, C++ and Fortran programs that use the POSIX pthreads threading primitives.
 - ▶ Unlocking an invalid and not-locked mutex
 - ▶ Destroying an invalid or a locked mutex
 - ▶ Invalid or duplicate initialization of a pthread barrier
 - ▶ Waiting on an uninitialized pthread barrier
 - ▶ Many others

THREAD ERROR DETECTORS:

▶ DRD

- ▶ Detecting errors in multithreaded C and C++ programs.
 - ▶ Data Race Detection
 - ▶ Lock contention. One thread blocks the progress of one or more other threads by holding a lock too long.
 - ▶ Deadlock. A deadlock occurs when two or more threads wait for each other indefinitely.
 - ▶ False sharing
 - ▶ Many others

CONCLUSIONS

WHY USE VALGRIND

- ▶ Valgrind will **save you hours** of debugging time
- ▶ Valgrind can help you **speed up your programs**
- ▶ Valgrind is **free**
- ▶ Valgrind is **widely used**
- ▶ Valgrind is **actively maintained**

WHEN SHOULD YOU USE VALGRIND?

- ▶ Easy answer: **all the time.**
 - ▶ After **big changes**
 - ▶ When a **bug** occurs (or is suspected)
 - ▶ **Before a release**
 - ▶ Whenever you want information about how your **program is spending its time**, or you want to **speed it up**

END